

Commodify

Team Members

Name	Email	GitHub
Iris Tiong	iristyx@seas.upenn.edu	iristyx
Francis Featherby	ffrancis@seas.upenn.edu	ffeatherby
Peter Brice	pbrice@seas.upenn.edu	briceybrit
Steven Brooks	sbr@seas.upenn.edu	stevegbrooks

Introduction

Commodify is an app to visualize and interact with commodity market and weather data.

The problem

Existing sources for commodity data deal mostly with individual sectors such as energy or agriculture, but rarely with the whole commodities space, and often with little analysis or interpretation. To get an overview of the entire market, one needs to manually integrate data from many different resources, which is time intensive and technically difficult for most traders and analysts. Commodify is the solution.

The solution

Commodify provides a “one-stop-shop” for commodities and weather data, and aims to be a useful resource for trade houses, banks and hedge funds analysing and trading commodities - and the many other markets influenced by commodities.

Architecture

List of Technologies

Exploratory data analysis (EDA) was performed on the downloaded raw datasets to check data availability (presence or absence of null values), usefulness of data, units, and in general to summarize the data characteristics. Subsequently, only relevant data columns, identified to be useful to the project goal, were extracted from the full datasets. Python was used at this stage due to its easy-to-use libraries, including pandas and numpy, which are widely used in data science applications. Microsoft Excel was also used for organizing some data. The processed datasets were hosted on AWS using MySQL.

The web application was developed using React, a user interface framework developed by Facebook. Using the Node package manager, a boilerplate application was set-up via the Node package create-react-app. A list of dependencies (required Node modules) for the client is outlined in the json file “commodify/app/client/package.json”. The professional look-and-feel and responsiveness of the webpage was implemented using Bootstrap, an open-source CSS framework, which contains built-in templates for interface components such as navbars and grid layouts.

On the backend we used Express and Node to write handlers for requests and to connect the web application with the database, so as to return results from user queries. A list of dependencies (required Node modules) for the server is outlined in the json file “commodify/app/server/package.json”.

Markdown was used to write the final report and other project documents. Markdown is a lightweight markup language used for creating formatted text using a plain-text editor.

Description of System Architecture

Commodity allows the user to search for raw data on commodity supply and demand and also presents visualisations of that data. Furthermore, since weather and climate are crucial determinants of the supply and demand of many commodities, it returns weather data and charts relevant to the commodity search. This illustrates not only trends in the commodities markets, but also the interaction between weather and climate. The site has a homepage and two functional pages, a “Dashboard” and a “Search” page, as well as a static “Contact” page designed for Commodity users to get in touch with the development team for any enquires, complaints, or suggestions.

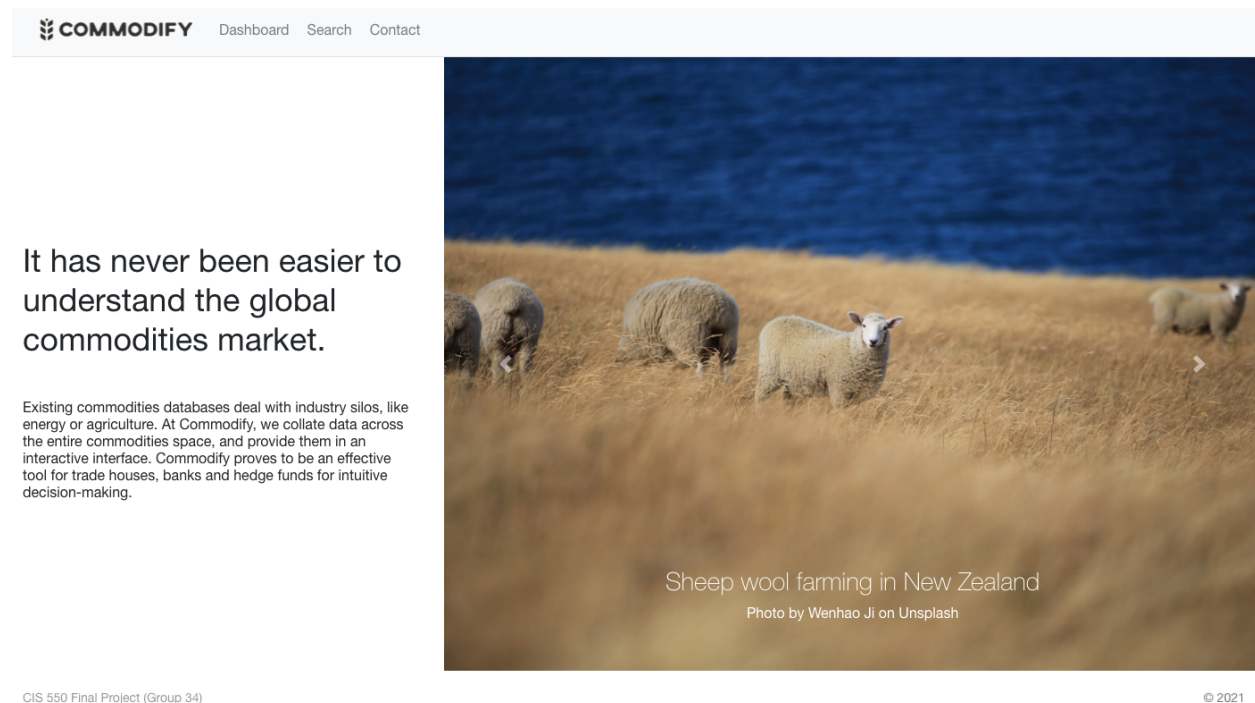


Figure 1:


Dashboard

The dashboard page gives the user a overview of the commodities market in graphical display formats. For the purpose of the project, the information presented on this page is for the USA only, due to the fact the data is most complete and available for the USA. The page contains two charts: the first chart presents a broad overview of production levels of commodities (soy, corn and wheat are the major stocks in USA) and electricity. Total rainfall level is overlaid and presented on the same chart to illustrate the broad relationship between rainfall levels and commodity production levels, if any. The second chart presents the total electricity production of each of the 50 USA states, levels illustrated by the colors (with red being the highest, and light yellow being the lowest).

Search

The search page allows the user to select a commodity sector, for example ‘agriculture’, or ‘renewables’. After selecting one of these options the user will be presented with the option to search for country or USA state level data, and will then be shown drop-down menus for the data available. Once the selections are made and the ‘Submit’ button is pressed, the user will see the commodity’s production and consumption data in graphical display (line chart) and tabular form. Should the user select to view commodity data for a USA

Trends in Commodities, Electricity, and Rainfall since 1990 in the United States



Electricity Production by State




Figure 2:

state (instead of country), an additional query is performed to retrieve weather information for the selected state. Monthly average rainfall and temperature data are then presented in a separate chart and tabular format below the first chart.

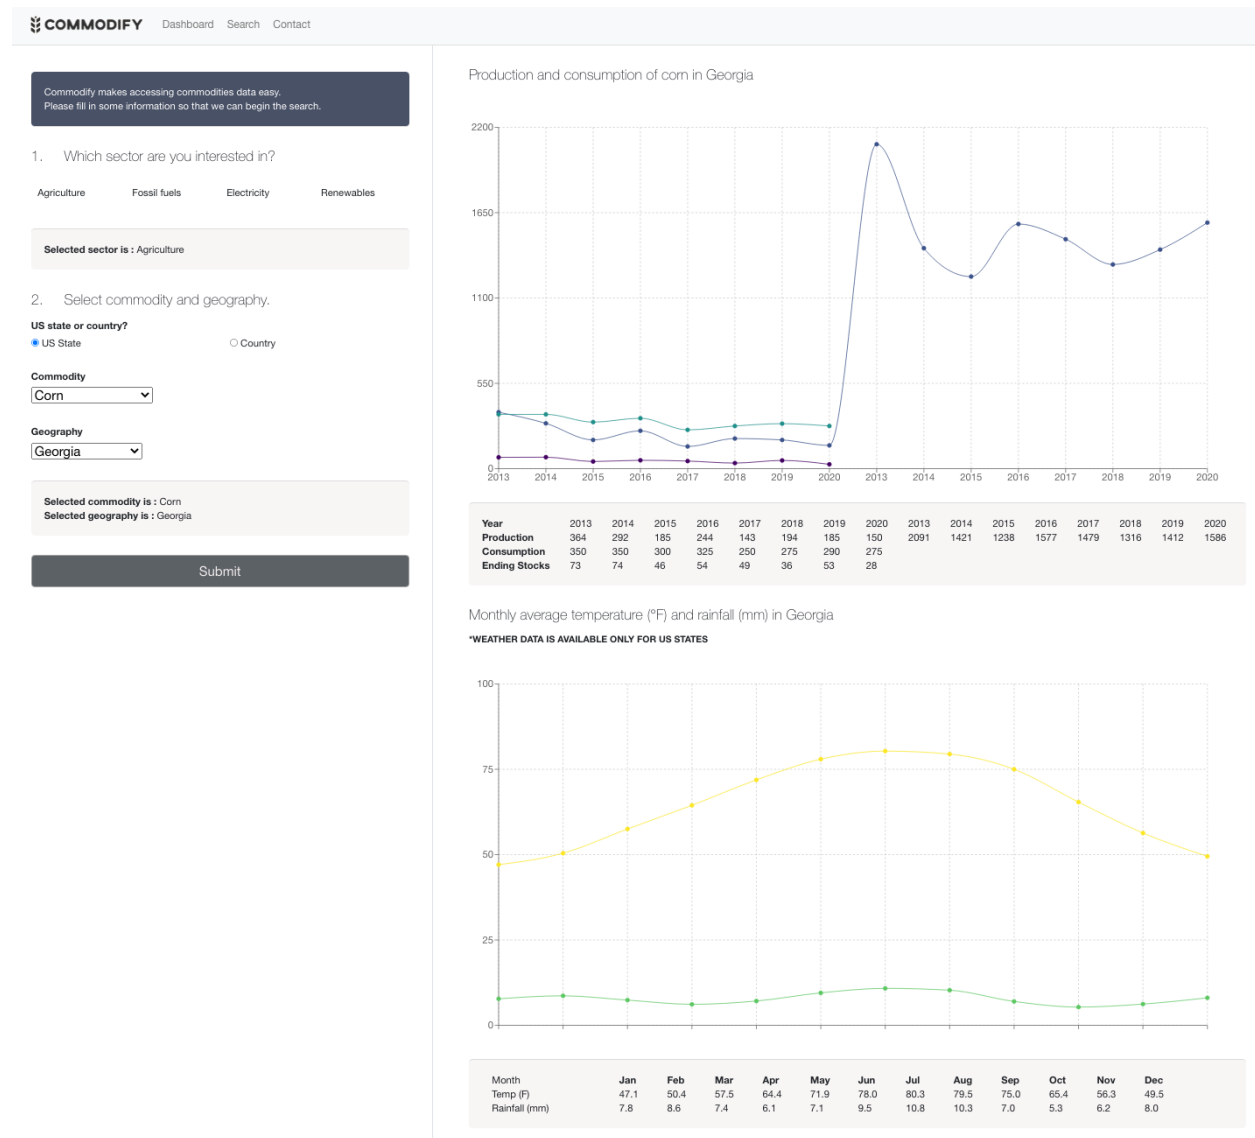


Figure 3:

Data

The datasets used in the application mainly agricultural commodities data, energy supply and demand data, as well as weather data obtained from public domains.

1. Current and historical agricultural commodities data from the USDA Foreign Agricultural Service

- Description: a dataset containing commodities and their prices around the world, including trading value at the beginning and end of each month.

- Size: ~200 MB with ~2 million rows and 12 features.
- A commodity is identified by “Commodity_Code”, and also has location, time, and various attributes attached to its value.
- A quick peek at the first 5 rows of the data shows the following (column names changed for brevity; not all columns shown):

cocode	comm	cntry	mrkt_yr	year	month	attr_id	attr	unit_id	unit_desc	val
577400	Almonds	AF	2010	2018	10	20	Beginning Stocks	21	(MT)	0
577400	Almonds	AF	2010	2018	10	125	Domestic Consumption	21	(MT)	0
577400	Almonds	AF	2010	2018	10	176	Ending Stocks	21	(MT)	0
577400	Almonds	AF	2010	2018	10	88	Exports	21	(MT)	0
577400	Almonds	AF	2010	2018	10	57	Imports	21	(MT)	0

- Number of unique values for each column (original column names):

col_name	unique
Commodity_Code	63
Commodity_Description	63
Country_Code	212
Country_Name	213
Market_Year	62
Calendar_Year	62
Month	13
Attribute_ID	71
Attribute_Description	71
Unit_ID	11
Unit_Description	11
Value	44725

- Summary statistics for numeric columns:

Market_Year	Calendar_Year	Unit_ID	Value
Min. :1960	Min. :1959	Min. : 2.000	Min. : -9510
1st Qu.:1979	1st Qu.:2006	1st Qu.: 8.000	1st Qu.: 0
Median :1992	Median :2006	Median : 8.000	Median : 15
Mean :1992	Mean :2006	Mean : 9.833	Mean : 11937
3rd Qu.:2007	3rd Qu.:2014	3rd Qu.: 8.000	3rd Qu.: 212
Max. :2021	Max. :2021	Max. :29.000	Max. :42528700

2. For energy and electricity data: Current and historical energy supply and demand data from the Energy Information Agency, part of the United States Department of Energy, Current and historical electricity supply and demand data from the Energy Information Agency, part of the United States Department of Energy, and Historical energy data from the Statistical Review of World Energy by BP

- Description: The EIA produces data sets released weekly and monthly which contain thousands of data points on US and global energy production and consumption. Two of the csv datasets which were used were: “Retail Sales of Electricity by State by Sector by Provider (EIA-861)” (analogous

to consumption of electricity) and “Net Generation by State by Type of Producer by Energy Source (EIA-906, EIA-920, and EIA-923).” Total electricity consumption from “Sales to Ultimate Customers” in megawatthours by state by year was manually merged with total (from all energy sources) electricity production/generation into a single csv.

- In addition, BP produces an annual report which contains a consolidated dataset, panel format, of annual production and consumption data of energy types for countries from 1965 to 2019. All values are in exajoule (EJ) except for crude oil which was converted from megatonnes using a conversion factor of 23.44(m)tonnes/EJ. This dataset was added to provide more country specific energy data, specifically data of consumption and production of Crude Oil, Coal, Natural gas and Renewable energy.
- Size: the data are provided in various packages and formats, some overlapping, but there are at least hundreds of features per week in data stretching back decades, i.e. thousands of rows.
- The data are mostly time series. They show various aspects of supply and demand, e.g. production, consumption, inventories etc., for a given geographical region, such as a state or country, on a sequence of dates of in a sequence of periods. For the final project, only production and consumption data by political identity, by year was used.
- Below are some example lines from a CSV file dealing with US crude oil production in the lower 48 United States by month in 2020 in thousands of barrels (column names changed for brevity).

Date	U.S. Prod	East Coast Prod	Florida Prod	New York Prod
2020-06-15	313264	1967	69	23
2020-07-15	340152	1968	122	23
2020-08-15	328099	2189	122	23
2020-09-15	326114	2309	108	23
2020-10-15	323387	2359	113	23
2020-11-15	333721	2180	117	23

- Summary statistics:

Date	US Prod	East Cost Prod	Florida Prod	New York Prod
Min. :1981-01-15 00:00:00	Min. :119208	Min. : 399.0	Min. : 35.0	Min. : 9.00
1st Qu.:1990-12-30 12:00:00	1st Qu.:175752	1st Qu.: 649.5	1st Qu.: 174.5	1st Qu.:19.00
Median :2000-12-15 00:00:00	Median :212585	Median : 861.0	Median : 370.0	Median :28.00
Mean :2000-12-14 11:43:27	Mean :223392	Mean :1171.7	Mean : 511.4	Mean :33.37
3rd Qu.:2010-11-30 00:00:00	3rd Qu.:263536	3rd Qu.:1518.0	3rd Qu.: 542.5	3rd Qu.:35.50
Max. :2020-11-15 00:00:00	Max. :396865	Max. :4243.0	Max. :3606.0	Max. :96.00

3. Current and historical weather data from the US National Oceanic and Atmospheric Administration (NOAA) National Center for Environmental Information (NCEI)

- Description: a dataset containing daily weather data for weather stations in the USA, including, from 1929 to the present (2021).
- Size: ~ over 20 GB with 28 features, including temperature, dewpoint, surface elevation, rainfall, wind speed, etc. The full dataset was downloaded and relevant information was extracted. Only state, rainfall and temperature data were retained for simplicity and usefulness.
- A quick peek at the first 5 rows of the data shows the following (column names changed for brevity; not all columns shown):

Entity_id	Year	Month	Temperature	Precipitation
108	1931	1	30.35	70.96
108	1931	2	35.96	46.42
108	1931	3	43.77	35.48.
108	1931	4	10.00	10.00.
108	1931	5	12.90	12.90

- Number of unique values for each column (selected useful column):

col_name	unique
Entity_id	53
Year	91
Month	12
Temperature	7524
Precipitation	5985

- Summary statistics for numeric columns:

Entity_id	Year	Temperature	Precipitation
Min. :1	Min. :1931	Min. :-10.2	Min. :0
1st Qu.:109	1st Qu.:1960	1st Qu.:41.06	1st Qu.:2.9
Median :159	Median :1981	Median :56.54	Median :10.58
Mean :152	Mean :1980	Mean :54.82	Mean :17.84
3rd Qu.:197	3rd Qu.:2001	3rd Qu.:70.25	3rd Qu.:30.35
Max. :262	Max. :2021	Max. :94.20	Max. :87.55

Database

The commodities and weather data were sourced as txt and csv files and processed using python notebooks such as commodity_usda_state.ipynb in our submission. This removed problem values, made names and units consistent, dropped unwanted columns and so on. In particular, the massive weather dataset (over 20GB of csv files) provided information in terms of daily statistics collected across virtually all global weather stations, from the years 1929 to the present (2021). In order to render the dataset useful for the application, data aggregation was performed to obtain average monthly statistics in US states only.

The processed data were output as csv files and then uploaded to the database via MySQL. The database can be recreated using the DDL.sql file which links to the provided csv files.

The major entity resolution questions concerned how to deal with the problem that different commodities have different attributes (for example corn has acreage but electricity does not), and different data are available at the country and state levels, for example wheat production data are available by country and by US state, but wheat consumption is only available at the country level. We decided that the best solution for the user was to include all the commodities supply and demand data in a single table, called **Commodity**, meaning we avoided having many different tables with different attributes at the cost of having many null values in the table. For example, Commodity was not split into Agriculture and Non-agriculture tables even though only production and consumption data is populated for Non-agriculture commodities (the rest of the non-primary key values are null).

The **Commodity** table is in BCNF. We achieved this by placing metadata, such as the sector each commodity belongs to (e.g. agriculture for soybeans), to a separate table called **Commodity_Group**, and also the

Political_Entity table, which contains information about countries and US states, including their names and the id used in the **Commodity** table. (It was necessary to use ids rather than names because there is a country called Georgia and also a state!) We determined which territories to include and which versions of their names to use, which abbreviations, and their ids, manually in Excel. The **Commodity_Group** and **Political_Entity** tables are also in BCNF. It was necessary to decide which countries and territories to include (for example we excluded the Netherlands Antilles from our list of political entities, although we had historic data for them, because that country was dissolved in 2010), which exact names to use for them, and which abbreviations.

The weather data are held in a table called **Weather**, which is also in BCNF. While a range of climate parameters were included in the raw dataset, the team identified that rainfall and precipitation were the most relevant weather parameters to study commodities trends, and hence, other weather data was omitted (e.g. wind speeds, haze) at this stage of the project to avoid unnecessary complexity of the dataset.

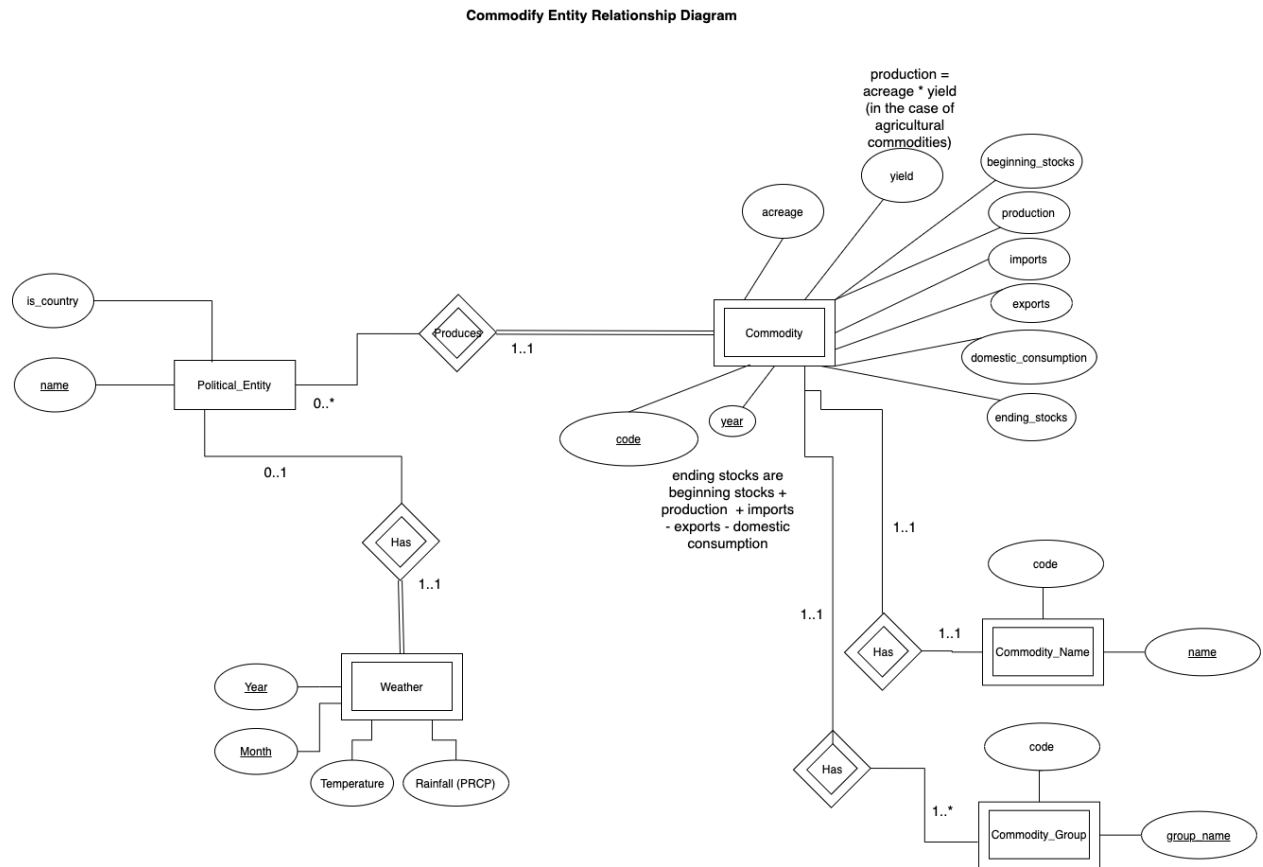


Figure 4:

Relation schema:

- Commodity (name, year, month, pe_id, beginning_stocks, ending_stocks, imports, exports, acreage, yield, production, consumption)

Total instances in Commodity: 177,457

- Commodity_Group (name, group_name)

Total instances in Commodity_Group: 68

- Political_Entity (id, name, is_country, abbrev, geo_id)

Total instances in Political_Entity: 269

- Weather (pe_id, year, month, temp, rainfall)

Total instances in Weather: 50,136

Political Entity (pe_id) provides the foreign key between Commodity and Weather.

Queries

Building the search function involved devising numerous complex searches to ensure that at each stage of the selection process the user sees only options which will result in non-null results. For example, the following query ensures that once the user has selected a commodity sector and whether they want state or country-level data, they will only be able to choose from commodities for which the database holds data for some of the relevant attributes in 2019 (the last year with full data available):

```
SELECT DISTINCT C.name
FROM Commodity C JOIN Political_Entity P ON C.pe_id=P.id JOIN Commodity_Group G ON C.name = G.name
WHERE P.is_country='${eT}' AND C.year=2019 AND G.group_name='${sector}' AND ((C.production != 0 OR C.consumption != 0 OR C.ending_stocks != 0)
OR (C.consumption != 0 OR C.consumption != null) OR (C.ending_stocks != 0 OR C.ending_stocks != null))
ORDER BY C.name ASC;
```

Likewise, the below query returns a list of only those states or countries for which there is data for 2019 for one of the relevant attributes for the commodity previously selected:

```
SELECT name
FROM Political_Entity
WHERE is_country = '${eT}' and name IN (
    SELECT DISTINCT P.name
    FROM Commodity C JOIN Political_Entity P ON C.pe_id=P.id
    WHERE C.name = '${commodity}' AND ((C.production != 0 OR C.consumption != null)
    OR (C.consumption != 0 OR C.consumption != null) OR (C.ending_stocks != 0 OR C.ending_stocks != null))
```

Once a commodity and state or country have been selected, the below query returns the relevant production, consumption and ending stock data for recent years:

```
SELECT C.year, C.production, C.consumption, C.ending_stocks
FROM Commodity C JOIN Political_Entity P on C.pe_id = P.id
WHERE C.name = '${commodity}' and P.name = '${entity}' AND C.year > 2012;
```

Finally, if the user selected a state, the below query returns climate data for that state, specifically average temperature and rainfall for each month, based on data for the last fifty years:

```
SELECT month, AVG(temp) AS temp, AVG(rainfall) AS rainfall
FROM Weather W JOIN Political_Entity P ON W.pe_id=P.id
WHERE P.name = '${state}' and W.Year > 1970
GROUP BY month
ORDER BY month ASC;
```

The above queries completed in a satisfactory time.

The Dashboard page used the following queries to obtain the data needed for the chart:

```
SELECT C.year,
SUM(IF(C.name='Oilseed, Soybean', production, NULL )) AS soy_prod,
SUM(IF(C.name='Corn', production, NULL )) AS corn_prod,
SUM(IF(C.name='Wheat', production, NULL )) AS wheat_prod,
ROUND(SUM(IF(C.name='Electricity', production, NULL ))/100000) AS elec_prod,
```

```

ROUND(SUM(W.rainfall)) AS rainfall
FROM commodify.Commodity C
  JOIN commodify.Political_Entity PE ON C.pe_id = PE.id
  JOIN commodify.Weather W ON PE.id = W.pe_id
WHERE C.year < 2020 AND C.year >= 1990 AND PE.is_country = 0
GROUP BY C.year;

```

Benchmarking

We used the following to assess the performance of the original query:

```

set profiling=1;
FLUSH STATUS;
SELECT sql_no_cache ORIGINAL_QUERY;
SHOW profile;

```

Which showed us the following (average of 10 runs):

Status	Duration
Opening tables	0.000616
init	0.000008
System lock	0.000001
optimizing	0.000019
statistics	0.000061
preparing	0.000031
Creating tmp table	0.000042
executing	4.651374
end	0.000013
query end	0.000005

Clearly, the most expensive thing is the actual execution.

To look at number of page I/Os, we used:

```

FLUSH STATUS;
SELECT sql_no_cache ORIGINAL_QUERY;
SHOW STATUS LIKE 'last_query_cost';

```

Which showed us that there are around 200,000 page reads for this query.

Optimization

- Instead of using `Commodity.year` for the selection and grouping, we used `Weather.year`.
 - `Commodity` is a much bigger table: about 175,000 rows vs. 50,000 rows.
 - This reduced the page reads by half, but did not affect the query speed.
- We then investigated further and found that we only have `Commodity` data *by year* in the US, but we have `Weather` data *by month*. So when the JOIN happens, it basically does a cartesian product of the cardinality of `Commodity` times the cardinality of `Weather`. The solution then became obvious - use a subquery for `Weather` that groups it by year before the the JOIN:

```

...
SELECT W.year,
  SUM(IF(C.name = 'Oilseed, Soybean', C.production, NULL )) AS soy_prod,
  SUM(IF(C.name='Corn', C.production, NULL )) AS corn_prod,

```

```

SUM(IF(C.name='Wheat', C.production, NULL )) AS wheat_prod,
ROUND(SUM(IF(C.name='Electricity', C.production, NULL ))/100000) AS elec_prod,
W.rainfall AS rainfall
FROM (SELECT * FROM (
    SELECT year, ROUND(SUM(rainfall)) AS rainfall
    FROM commodify.Weather
    GROUP BY year
    ) AS wby
) AS W
JOIN commodify.Commodity C ON W.year = C.year
JOIN commodify.Political_Entity PE ON PE.id = C.pe_id
WHERE PE.is_country = 0
GROUP BY C.year;
```

```

- The page reads fell to just over 5,000, and the query time fell to around half a second - a 10x improvement.

Below is the profile table after making the optimization (average of 10 runs):

| Status             | Duration |
|--------------------|----------|
| Opening tables     | 0.000091 |
| init               | 0.000007 |
| System lock        | 0.000009 |
| optimizing         | 0.000004 |
| statistics         | 0.000017 |
| preparing          | 0.000010 |
| Creating tmp table | 0.000029 |
| executing          | 0.055620 |
| end                | 0.000012 |
| query end          | 0.000005 |

## Technical challenges

Perhaps the greatest challenge was formatting the website, including building the .css style files and getting user interfaces such as buttons and dropdown menus to function correctly. We struggled, for example, to get data to render in rows rather than a single column. Without much experience of web design or any obvious single resource for guidance we resorted to a lot of googling (often unsuccessfully) and trial and error.

## Appendix

### Instructions to Build Locally

#### 1. Setup Connection to Database

Follow these steps for connecting to “commodity-db” on AWS with READ ONLY access:

Username: guest Password: wheat

Port: 3306

Endpoint: commodity-db.cn5sga8k6aq8.us-east-1.rds.amazonaws.com

Run this line in your terminal:

```
mysql --host=commodify-db.cn5sga8k6aq8.us-east-1.rds.amazonaws.com --port=3306 --user=guest -p
```

Input password at prompt.

## 2. Launch client and server

In the project directory, you can run:

### **npm start**

Runs the app in the development mode.

Open <http://localhost:3000> to view it in the browser.

The page will reload if you make edits.

You will also see any lint errors in the console.

### **npm test**

Launches the test runner in the interactive watch mode.

See the section about running tests for more information.

### **npm run build**

Builds the app for production to the **build** folder.

It correctly bundles React in production mode and optimizes the build for the best performance.

The build is minified and the filenames include the hashes.

Your app is ready to be deployed!

See the section about deployment for more information.

## List of Dependencies

- Backend: Node.js server with MySQL database hosted on AWS.
- Frontend: React.js with **react-chartist** for plots and charts and **react-leaflet** for maps.

## Client

```
"dependencies": {
 "@testing-library/jest-dom": "^5.11.9",
 "@testing-library/react": "^11.2.5",
 "@testing-library/user-event": "^12.7.1",
 "bootstrap": "^4.6.0",
 "chartist": "^0.11.4",
 "d3-fetch": "^2.0.0",
 "d3-scale": "^3.3.0",
 "d3-scale-chromatic": "^2.0.0",
 "prop-types": "^15.7.2",
 "raw-loader": "^4.0.2",
 "react": "^17.0.1",
 "react-chartist": "^0.14.4",
 "react-dom": "^17.0.2",
```

```
"react-markdown": "^6.0.1",
"react-router-dom": "^5.2.0",
"react-scripts": "^4.0.2",
"react-simple-maps": "^2.3.0",
"recharts": "^2.0.9",
"web-vitals": "^1.1.0"
},
```

## Server

```
"dependencies": {
 "body-parser": "^1.19.0",
 "cors": "^2.8.5",
 "express": "^4.17.1",
 "expresss": "^0.0.0",
 "mysql": "^2.18.1"
}
```